

Topics

- What is a Container Manager?
- What is Kubernetes?
- Kubernetes Core Concepts?
 - POD
 - Deployment
 - Services
 - Namespaces
- Kubernetes Advanced Concepts
 - Ingress
 - Secrets
 - Config Maps
 - Custom Resource Definitions

Container Bootcamp

Kubernetes Training

2: What is a Container Manager?




Cloud Native

**Applications adopting the principles of
Microservices packaged and delivered as
Containers orchestrated by
Kubernetes running on top of
Cloud infrastructure**

I,C,P as a Service: Terms

IaaS

Virtualization
Networking
ACLs
Firewalls
...



CaaS

Container Runtime
Images
Container Manager
...



PaaS

"Ready to Use" Solution Stack
Logging
Monitoring
Persistent volumes
...



Advantages *aaS

- **Cost Savings**
- **Pay what you need (although provider needs spare too)**
- **Faster Time to Market**
- **Scalability (for peak demands)**
- **Resiliency**
- **Easy business growth (or reduction)**

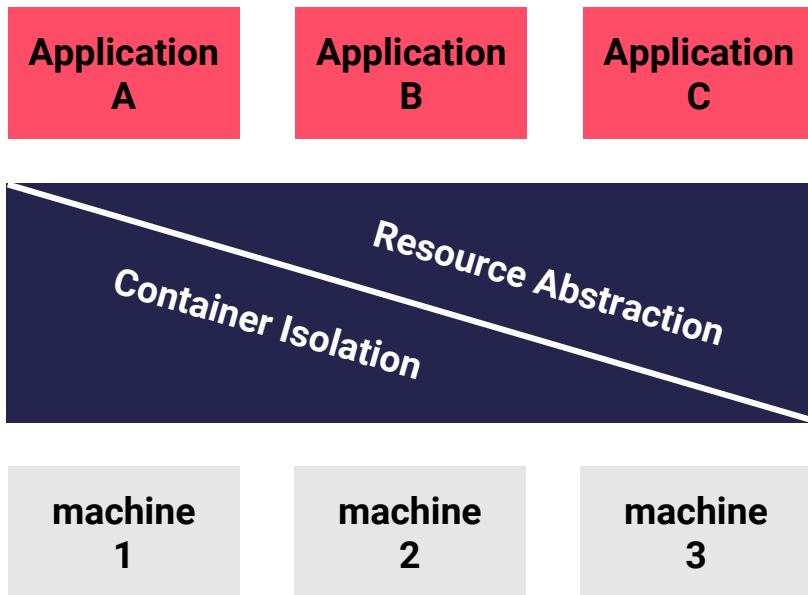
Is that all?

Without influence to

Development, Test, Deployment, Production, Architecture?

Aim: Ressourcen Abstraction

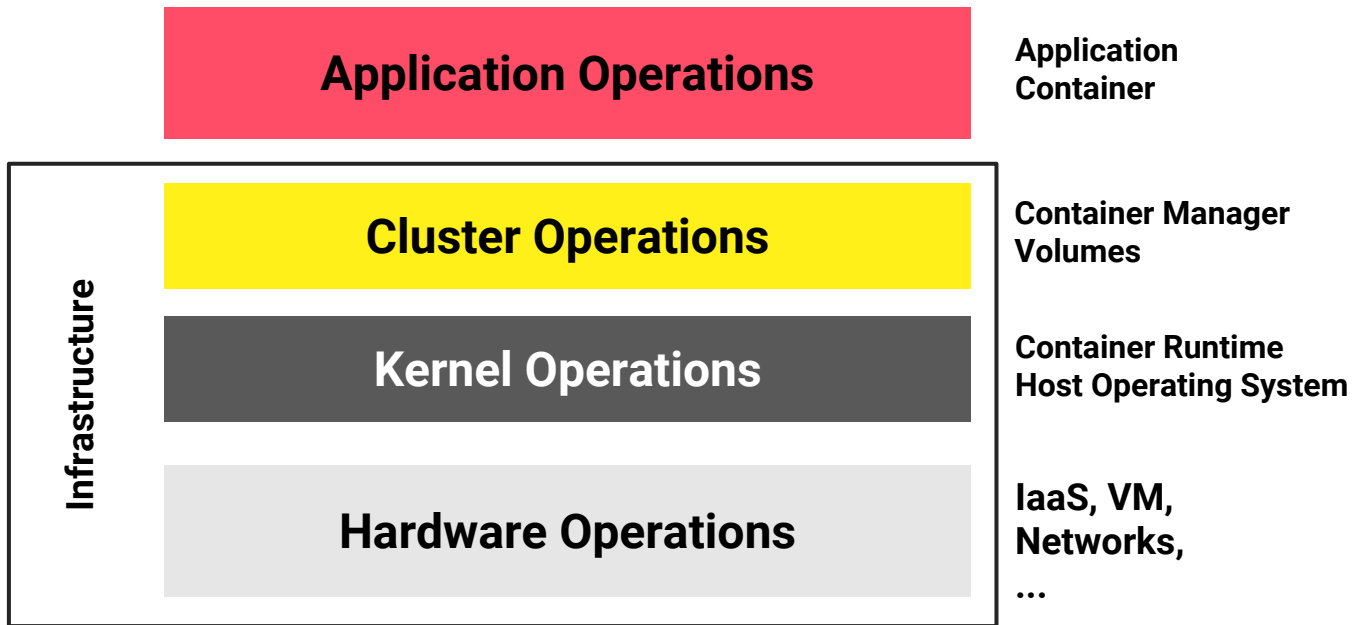
Applications consume Resources, not Machines



Why?

- Resiliency
- Scalability
- ...

Advantage: Decoupling Operations



Goals:

- 1. Any operations should be Application Oriented**
- 2. Platform should automate routine tasks**
(placement, healthchecks, healing, scaling, ...)
- 3. NFR should be moved to the platform**
(discovery, jobs, log aggregation, metrics collection, ...)
- 4. Allow developers to code the business domain**

Container Bootcamp

Kubernetes

3: What is Kubernetes?



What is Kubernetes?

Kubernetes is an open-source platform for

- **Automating deployment**
- **Scaling**
- **Operations of application containers across clusters of hosts**
- **Providing container-centric infrastructure.**

It is:

- **portable: public, private, hybrid, multi-cloud**
- **extensible: modular, pluggable, hookable, composable**
- **self-healing: auto-placement, auto-restart, auto-replication, auto-scaling**

Kubernetes is not


... an all-inclusive PaaS (Platform as a Service)

and does not:

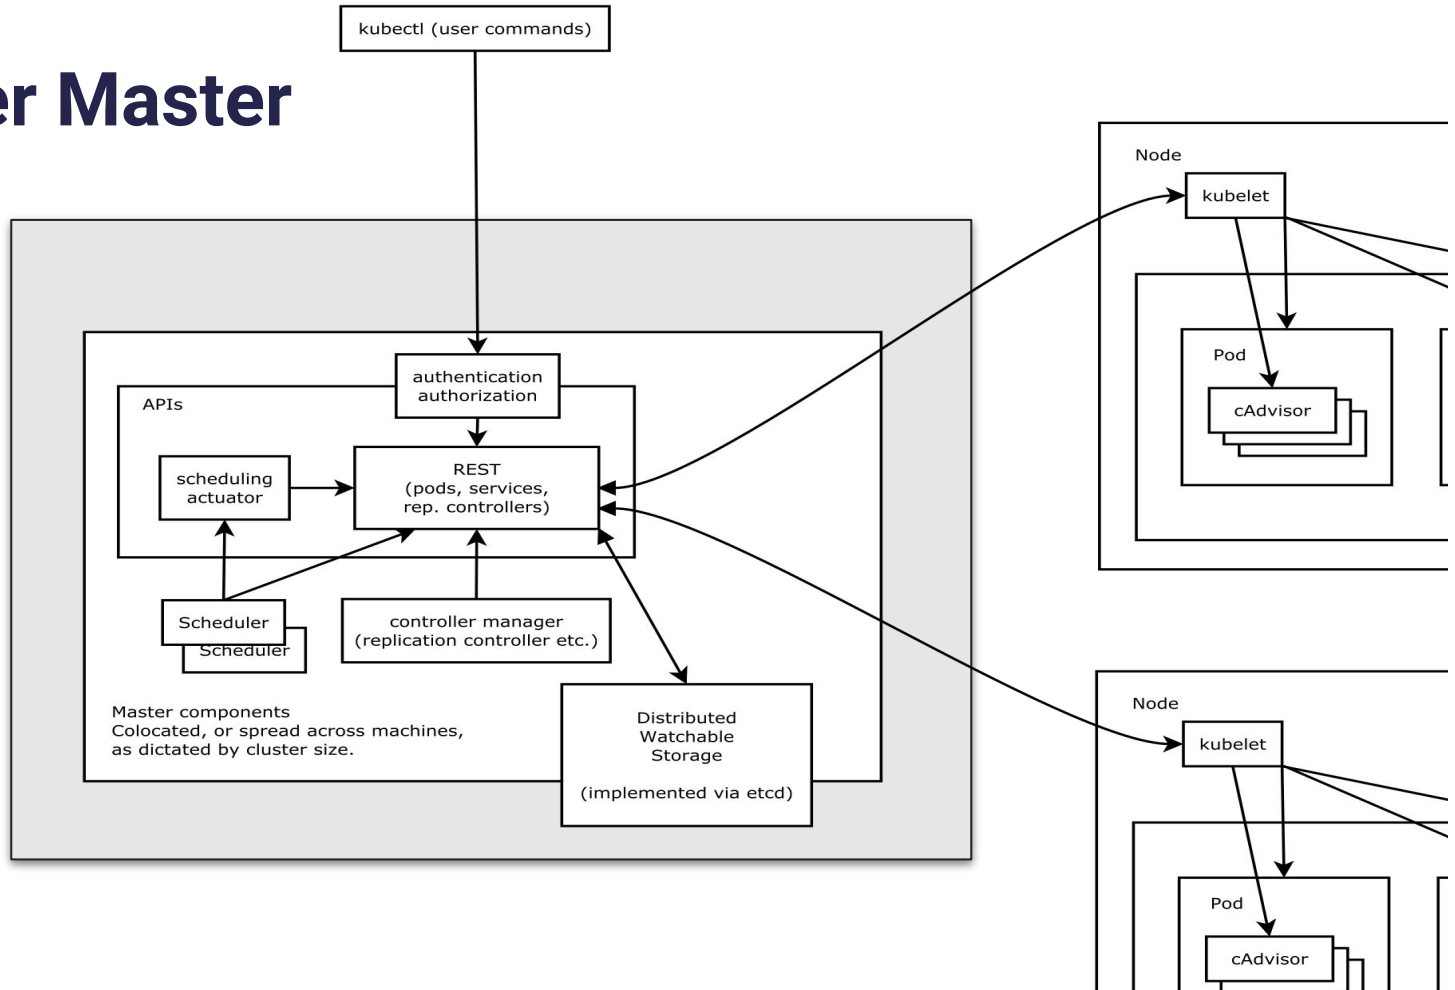
- **... dictate application frameworks**
- **... restrict the set of supported language runtimes**
- **... distinguish “Apps” from “Services”**
- **... provide any middleware or framework**
- **... provide a logging/monitoring/alerting or storage system**
- **... have an opinion in the source-to-image or CI/CD workflow space**
- **... have a click-to-deploy service marketplace**

K8s Processes

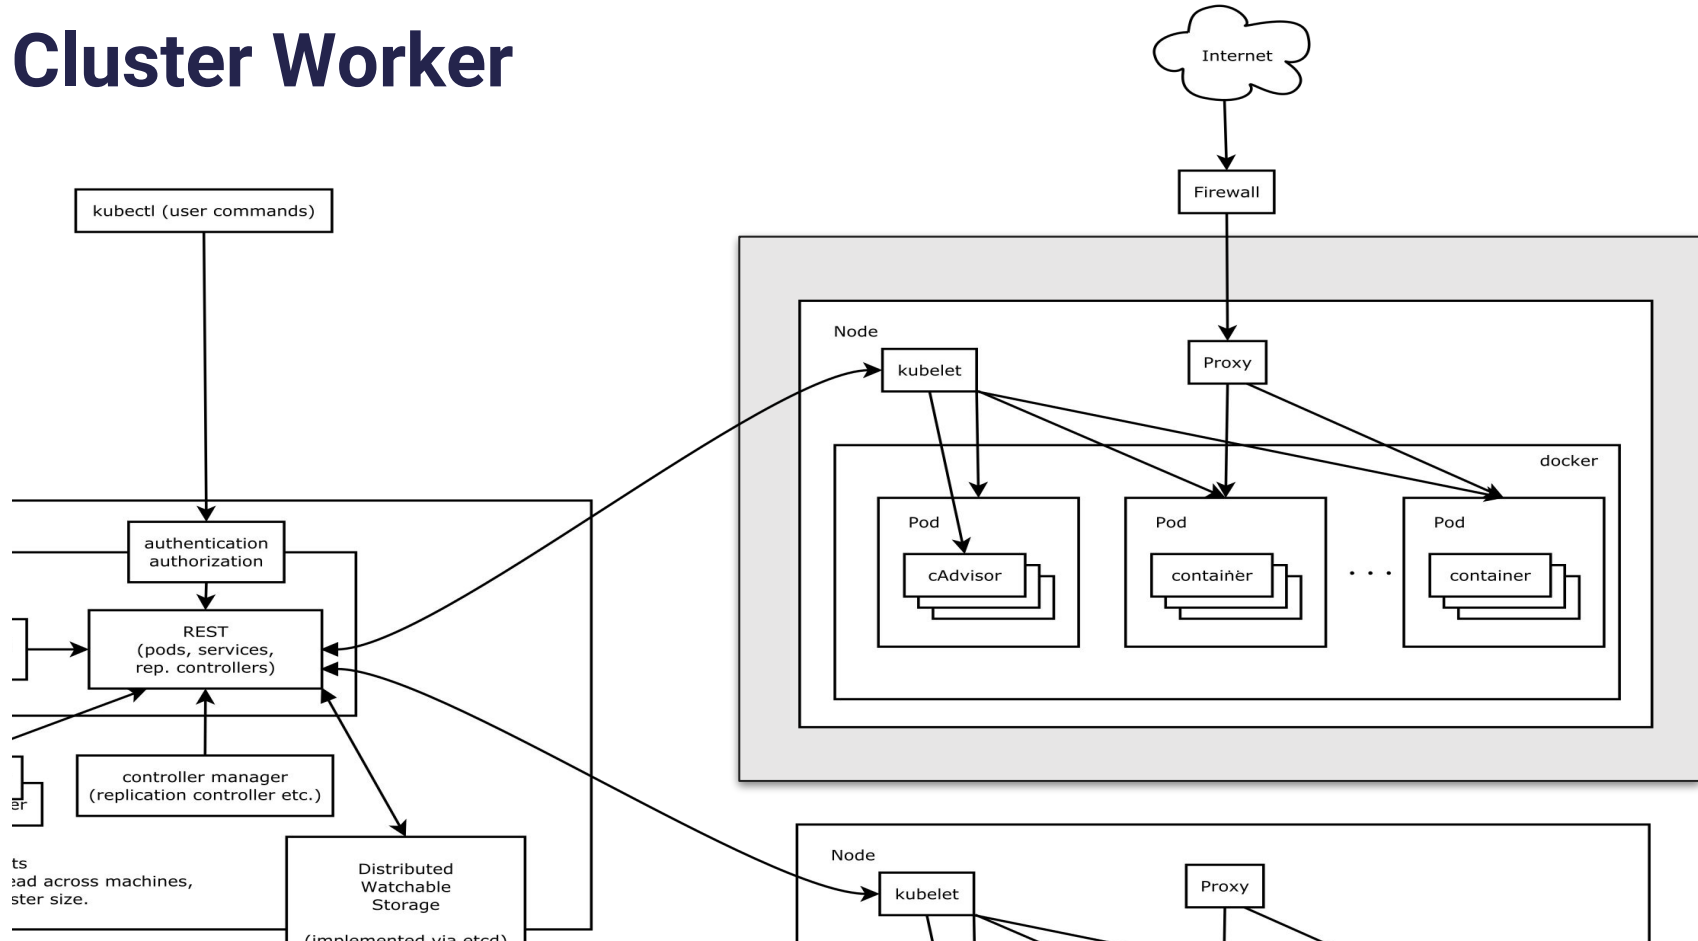
Cluster



Cluster Master



Cluster Worker



Cluster Modules

- **kubelet**
Manages pods and their containers, their images, their volumes, etc.
- **kube-proxy**
A distributed multi-tenant, round-robin load-balancer
- **etcd**
Distributed reliable key-value store, stores all persistent master state
- **API Server**
Mainly processes REST operations, validates them, and updates the corresponding objects in etcd
- **Scheduler**
binds unscheduled pods to nodes
- **Controller Manager Server**
Performs all other cluster-level functions (f.e. node discovery)

API Server

- **Is a HTTP API to your cluster**
 - For Users
 - For code running in containers
- **kubectl is the API client**
- **API reference (object reference) can be found here:**
<https://kubernetes.io/docs/api-reference/v1.9/>

...

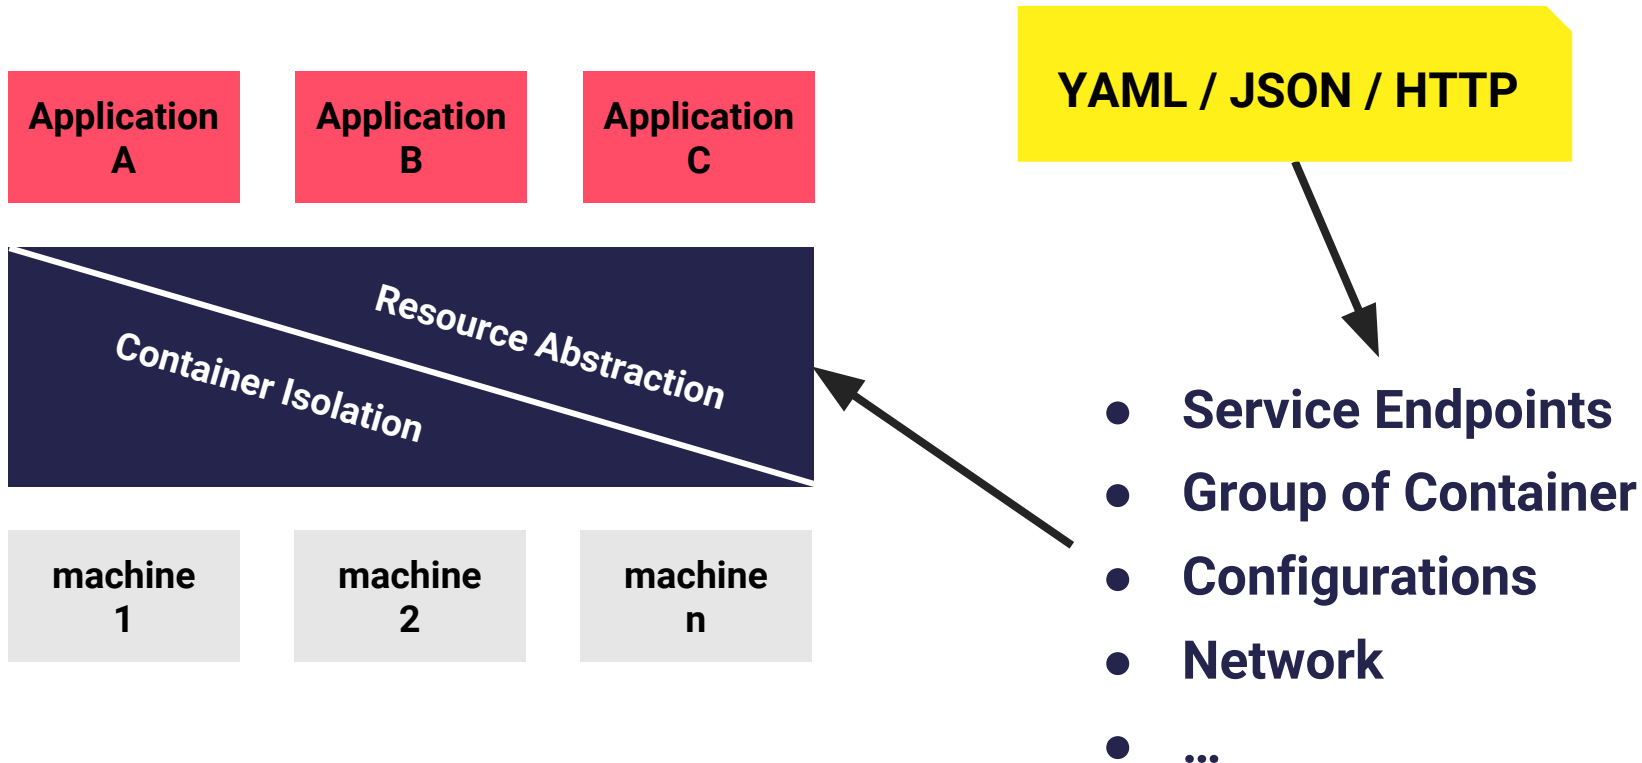
Container Bootcamp

Kubernetes

4: Kubernetes Core Concepts?



Ressourcen Abstraction



K8s Abstractions

- Pod
- Deployment
- Service / Service Discovery
- Persistent Volumes
- Config Maps
- StatefulSets
- Network Policies
- Ingress
- Images
- Labels and Selectors
- Replication Controller
- Secrets
- Names
- Namespaces
- Nodes
- Security Context
- Service Accounts
- Annotations
- Daemon Sets
- Ingress Resources
- Horizontal Pod Autoscaling
- Jobs
- Resource Quotas
- Replica Sets

Setup for Exercises

1. ~~<http://60minutek8s.ch.innoq.io/>~~ (daimler/tss)
2. Setup kubectl to point to our test cluster.
-> Copy the kubeconfig to a folder of your choice

```
$ export KUBECONFIG="$PWD/kubeconfig-X"  
$ set KUBECONFIG="/path/to/kubeconfig"  
$ kubectl get all  
No resources found.
```

3. Play around with `get pod`, `get node`, `get service` ...

Setup for Exercises

1. *Install Bash Completion*

```
$ source <(kubectl completion bash)
```

or

```
$ source <(kubectl completion zsh)
```

Bash Completion see: <https://kubernetes.io/docs/tasks/kubectl/install/>

Which Kubernetes Objects to Learn?

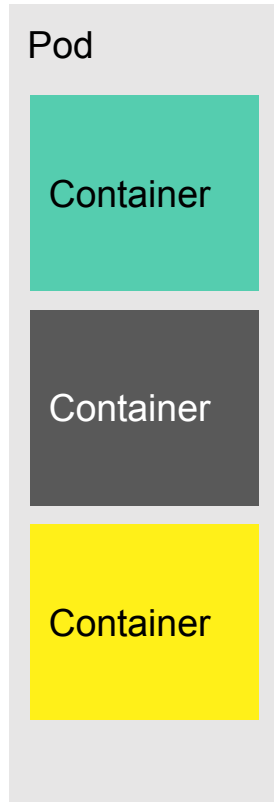
- **PODs**
- **Deployment**
- **Services**
- **Namespaces**

POD

What is a pod?


“Logical Host”

- it contains one or more containers which are tightly coupled
- co-located and co-scheduled
- run in a shared context
- shared storage (volumes)
- has its own IP



Why one or more container?

Sidecar / Sidekick pattern



POD example

```
apiVersion: v1
kind: Pod
metadata:
  name: redis-django
  labels:
    app: web
spec:
  containers:
    - name: key-value-store
      image: redis
    - name: frontend
      image: django
```

Exercise

1. Create a first Pod

```
$ watch kubectl get po  
  
$ kubectl create -f myclock.yaml
```

myclock.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: clock  
  labels:  
    app: training  
spec:  
  containers:  
    - name: my-clock  
      image: innoq/clock
```

```
$ kubectl get pod  
$ kubectl get po clock -o json  
$ kubectl get po clock -o jsonpath='{.spec.containers[0].image}'  
$ kubectl logs clock  
$ kubectl logs -f clock  
$ kubectl describe po clock
```

 JSONPath - XPath for JSON

Exercise

1. Create a Webserver Pod

```
$ kubectl create -f myserver.yaml  
  
$ kubectl port-forward server 8080:80
```

myserver.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: server  
  labels:  
    app: training  
spec:  
  containers:  
    - name: my-server  
      image: innoq/k8s-training-webserver
```

2. Open Terminal with `kubectl logs -f server`

3. Check <http://localhost:8080>

4. Check all known `kubectl` `get/describe` commands...

Can we connect to a Pod IP? Yes, we can...

```
$ kubectl describe po server
Name:          server
. . .
IP:           10.244.3.27
. . .
$ kubectl create -f busybox.yaml
$ kubectl exec busybox -it sh
/ # telnet 10.244.3.27 80
GET /
<!doctype html>
<html lang="en">
  <head>
    <title>K8s Training</title>
  </head>
  <body>
    <h1>Hello Kubernetes Training</h1>
  </body>
</html>
Connection closed by foreign host
/ # exit
```

```
apiVersion: v1
kind: Pod
metadata:
  name: busybox
spec:
  restartPolicy: Always
  containers:
  - image: busybox
    command:
      - sleep
      - "3600"
    imagePullPolicy: IfNotPresent
    name: busybox
```

Command, Args in a Pods definition

**When you create a Pod,
you can define a command
and arguments for the
containers that run in the
Pod.**

**This overwrites, whatever
is provided by the
container image**

```
apiVersion: v1
kind: Pod
metadata:
  name: command-demo
  labels:
    purpose: demonstrate-command
spec:
  containers:
    - name: command-demo-container
      image: debian
      command: ["printenv"]
      args: ["HOSTNAME", "KUBERNETES_PORT"]
```


A Pods Lifecycle

Lifecycle of a Pod

- **Pending:** The pod has been accepted by the system, but one or more of the container images has not been created
- **Running:** The pod has been bound to a node, and all of the containers have been created
- **Succeeded:** All containers in the pod have terminated in success, and will not be restarted
- **Failed:** All containers in the pod have terminated, at least one container has terminated in failure

Container Probes

Observable Interior



Container Probes

- **LivenessProbe:**

indicates whether the container is running

- On failure, container will be killed and subjected to its RestartPolicy
- Default state of Liveness before the initial delay is Success
- State when no probe is provided is Success.

- **ReadinessProbe:**

indicates whether the container is ready to service requests

- On failure, the pod's IP address will be removed from all services
- State before the initial delay is Failure
- State when no probe is provided is Success

Container Probes

A Probe is a diagnostic performed periodically on a container (!)

- **ExecAction:** executes a command inside the container (exit status code 0 on success)
- **TCPSocketAction:** a tcp check against the container's IP address (port is open on success)
- **HTTPGetAction:** an HTTP Get against the container's IP address (200 >= status code < 400, on success)

Exercise

1. Create a livenessProbe

```
$ kubectl delete po server  
  
$ kubectl create -f myserver.yaml
```

What happens?

What is CrashLoopBackOff?

How to fix this?

myserver.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: server  
  labels:  
    app: training  
spec:  
  containers:  
    - name: my-server  
      image: innoq/k8s-training-webserver  
      livenessProbe:  
        httpGet:  
          path: /healthz  
          port: 80  
        periodSeconds: 1
```

Example: Probe with HTTP

- when "host" is not defined, "PodIP" will be used
 - f.e. `host: my-host`
- when "scheme" is not defined, "HTTP" scheme will be used.

Only "HTTP" and "HTTPS" are allowed

- f.e. `scheme: HTTPS`

```
livenessProbe:
  httpGet:
    path: /healthz
    port: 8080
    httpHeaders:
      - name: X-Custom-Header
        value: Awesome
    initialDelaySeconds: 15
    timeoutSeconds: 1
```

Example: Probe with exec

- **Command is executed**
 - Initially after 5 seconds
 - Every 5 seconds

```
livenessProbe:  
  exec:  
    command:  
    - cat  
    - /tmp/healthy  
  initialDelaySeconds: 5  
  periodSeconds: 5
```

Example: Probe with exec

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    test: liveness
  name: liveness-exec
spec:
  containers:
    - name: liveness
      args:
        - /bin/sh
        - -c
        - touch /tmp/healthy; sleep 12; rm -rf /tmp/healthy; sleep 600
      image: busybox
      livenessProbe:
        exec:
          command:
            - cat
            - /tmp/healthy
          initialDelaySeconds: 5
          periodSeconds: 5
```


Container Probes: When to use what?

- **A container should crash on its own, whenever it encounters an issue**
 - > use `RestartPolicy (Always, OnFailure)`
 - > so no liveness probe needed
- **If you'd like to start sending traffic to a pod only when a probe succeeds**
 - > specify a `ReadinessProbe`
- **If a container wants the ability to take itself out of service for maintenance**
 - > specify a `ReadinessProbe`

Exercise

1. *Start pod* `liveness.yaml`

2. *Watch events with* `kubectl describe pod liveness-exec`


How fast does it get restarted?

And how often?

3. *Remove the Pod with* `kubectl delete pod liveness-exec`

Lifecycle Hooks

Lifecycle conformance



Deployment

Deployment

- A Deployment controller provides declarative updates for Pods and ReplicaSets, therefore
 - Supervises multiple pods across multiple nodes
 - Ensures that a specified number of pod “replicas” are running
 - Similar to a process supervisor
 - recommend, even if your application requires only a single pod

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
        role: backend
    spec:
      containers:
        - name: nginx
          image: nginx:1.12
          ports:
            - containerPort: 80
```

Deployment

- Deployment creates pods from template

Pod
Template
for Pod creation

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
        role: backend
    spec:
      containers:
        - name: nginx
          image: nginx:1.12
          ports:
            - containerPort: 80
```

Exercise

1. *Create a Deployment via nginx.yaml*
2. *What shows us `kubectl get all` ?*
3. *Describe one of it's created Pods and the Deployment (labels?)*
4. *Delete one `nginx-???` Pod. Does it get restarted?*
5. `kubectl delete deployment nginx --grace-period=0`


Services

What is a Service?

- **Provides Service Discovery by DNS**
- **Is a distributed multi-tenant, round-robin load-balancer**
- **is an abstraction which defines access to f.e. a logical set of Pods**

Service

- Is an object, similar to Pod
- Delegates ports
- Has a name
- Selects a set of endpoints
 - PODs as endpoints (using Selectors)
 - Specific endpoints (as Endpoints Object)
- “Speaks” UDP/TCP
- Has different *Service Types*



Exercise Service

1. Create a Service `www`:

```
$ kubectl create -f myservice.yaml  
  
$ kubectl describe service www
```

myservice.yaml

```
apiVersion: v1  
kind: Service  
metadata:  
  name: www  
spec:  
  ports:  
    - port: 80  
      protocol: TCP  
  selector:  
    app: nginx
```

2. Delete the pod, what happens to the service?

3. Create another `server2` pod, what happens to the service?

4. Execute: `kubectl exec busybox nslookup www`

Exercise Service

```
$ kubectl describe service www
```

```
Name:                www
Namespace:           chris
Labels:              <none>
Selector:            app=training
Type:                ClusterIP
IP:                  10.0.195.251
Port:                <unset> 80/TCP
Endpoints:           10.244.80.9:80,10.244.98.8:80
Session Affinity:    None
No events.
```

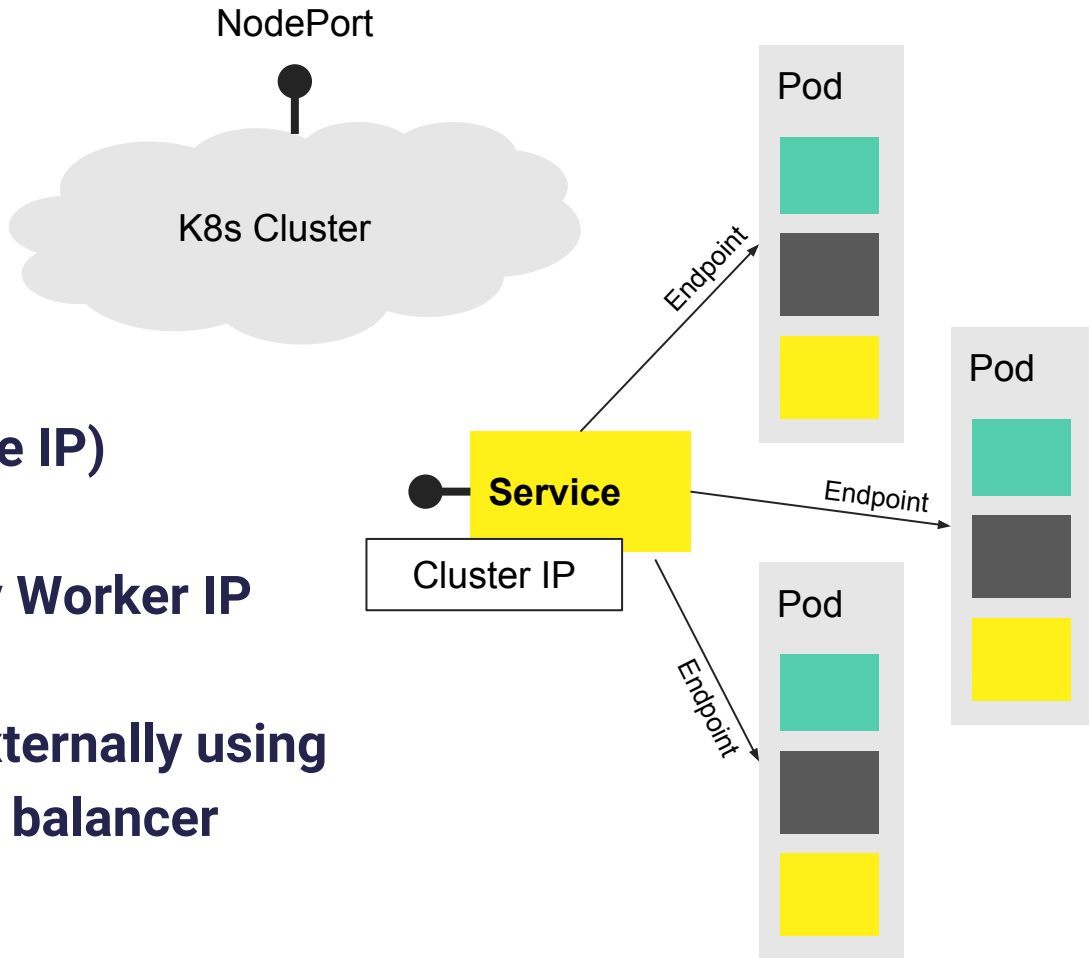
```
$ kubectl exec busybox nslookup www
```

```
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

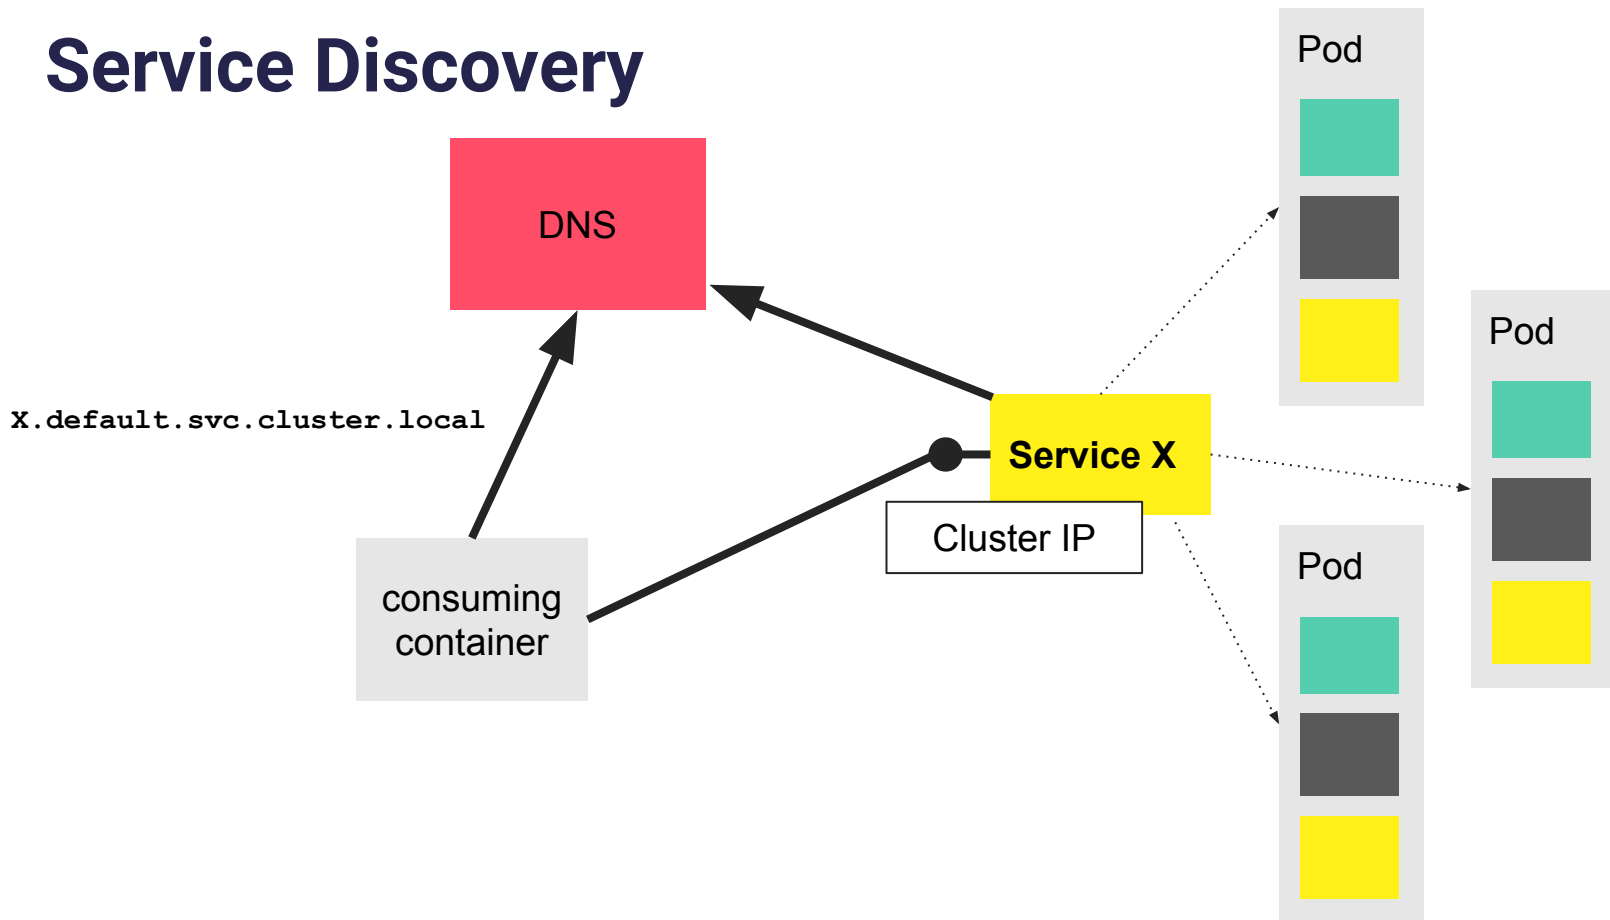
Name:      www
Address 1: 10.0.195.251 www.chris.svc.cluster.local
```

Service Types

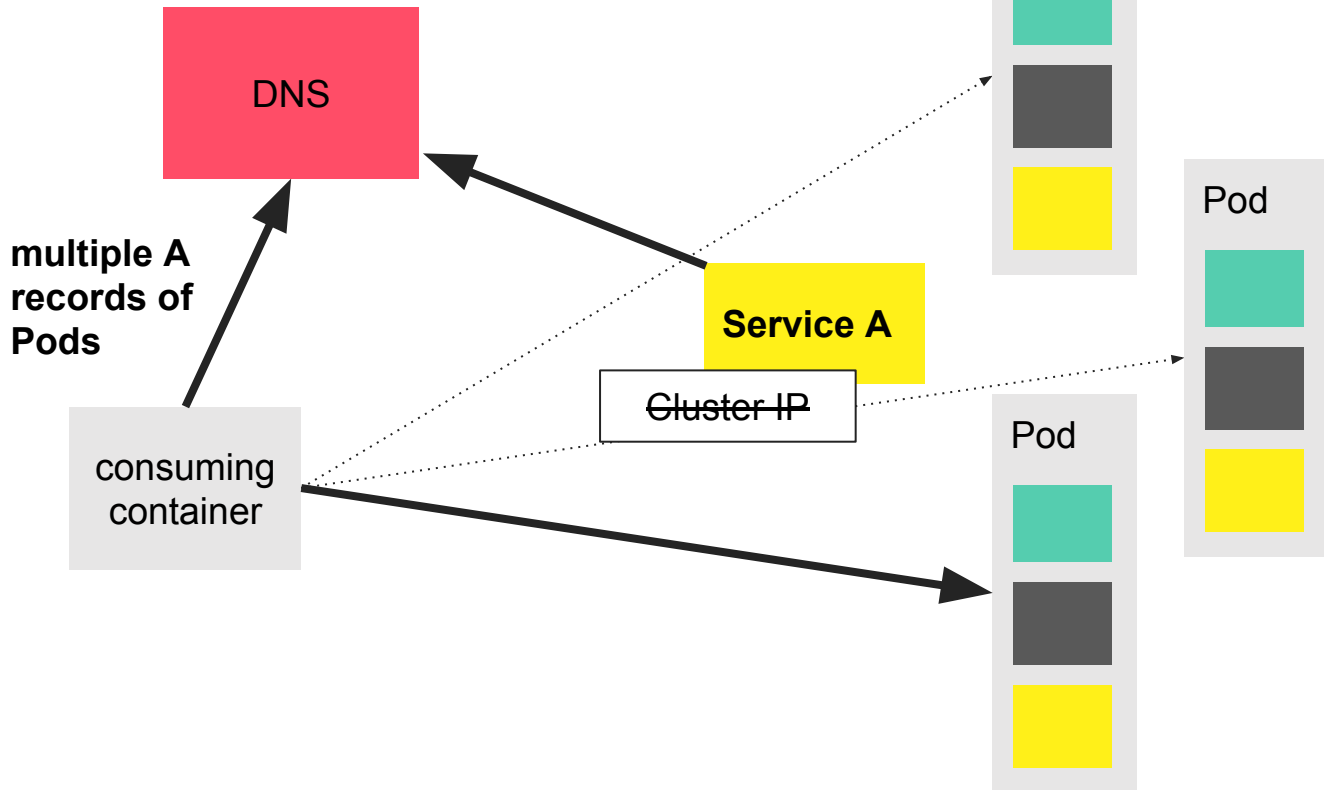
- **As a Cluster IP (Service IP)**
- **NodePort**
Port available on every Worker IP
- **Load Balancer**
Exposes the service externally using a cloud provider's load balancer



Service Discovery



Headless Service Discovery



```

apiVersion: v1
kind: Service
metadata:
  name: www
spec:
  clusterIP: "None"
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx
    
```

Exercise Service

```
$ kubectl exec busybox -i -t -- sh
/ # telnet wwwep 80
GET /
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html>. . .</html>
Connection closed by foreign host
/ # nslookup wwwext
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:        wwwext
Address 1: 2a02:2e0:3fe:1001:302:: redirector.heise.de
Address 2: 193.99.144.80 redirector.heise.de
/ # nslookup www
Server:      10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local

Name:        www
Address 1: 10.0.195.251 www.chris.svc.cluster.local
```


Namespaces

Namespaces

- **Are intended for use in environments with many users spread across multiple teams, or projects**
- **Provide a scope for names**
 - Names of resources need to be unique within a namespace, but not across namespaces
- **Are a way to divide cluster resources between multiple uses**
 - Resource quota
 - Namespace Isolation Policy (inbound traffic)
- **Service DNS entries have the form**
`<service-name>.<namespace-name>.svc.cluster.local`

Container Bootcamp

Kubernetes

5: Kubernetes Advanced Concepts



What to Learn?

- **Ingress**
- **Secrets**
- **Config Maps**

Ingress or How to Reach the Endpoints?

Ingress

Ingress is a collection of rules that allow inbound connections to reach the endpoints defined by a backend

Ingress can be configured to give services

- **externally-reachable urls and path**
- **load balance traffic**
- **terminate SSL**
- **offer name based virtual hosting**

Exercise Ingress

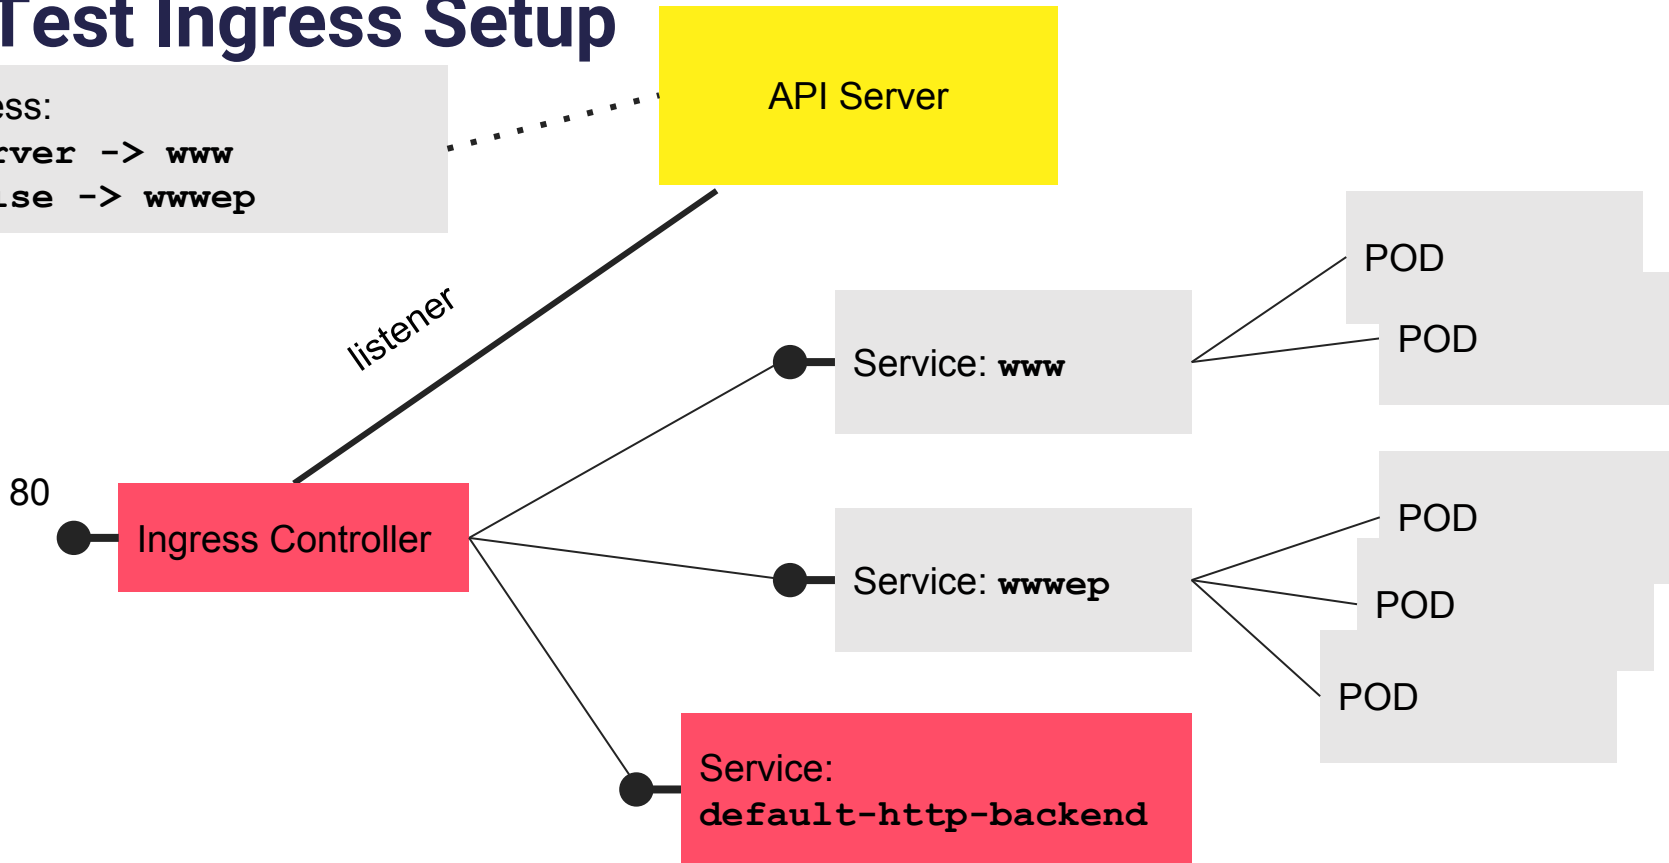
```
$ curl bootcamp.ch.innoq.io/heise -H 'Host: innoq.training.com'
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>301 Moved Permanently</title>
</head><body>
<h1>Moved Permanently</h1>
<p>The document has moved <a href="https://www.heise.de/">here</a>
</body></html>
$ curl 185.19.31.109/server -H 'Host: innoq.training.com'
<!doctype html>
<html lang="en">
  <head>
    <title>K8s Training</title>
  </head>
  <body>
    <h1>Hello Kubernetes Training</h1>
  </body>
</html>
```

myingress.yaml

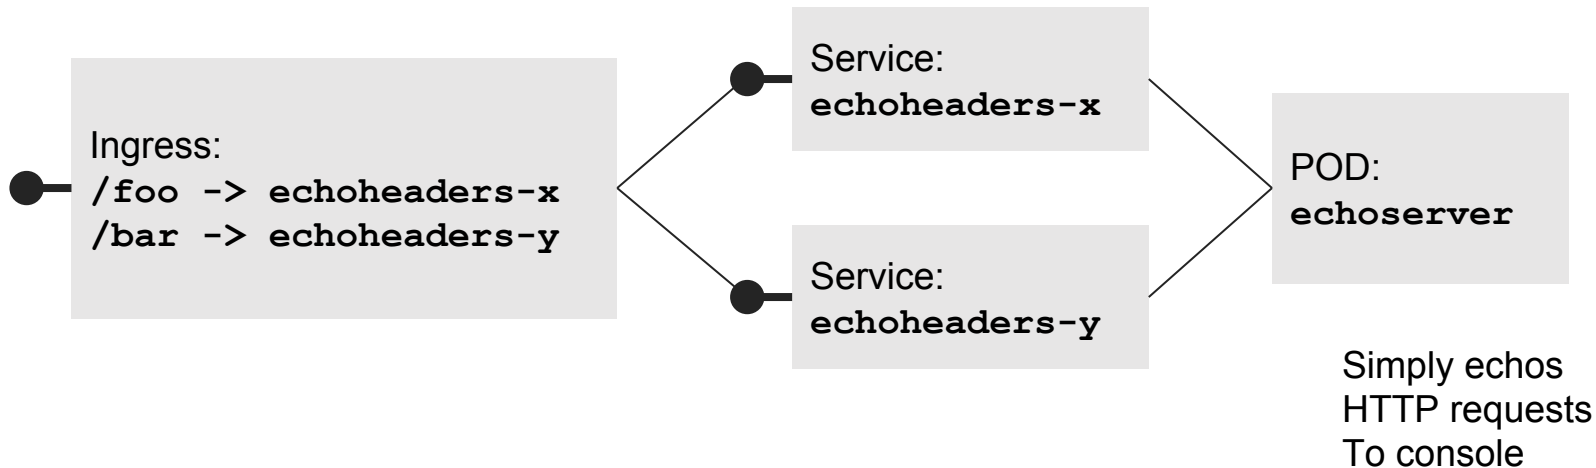
```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    ingress.kubernetes.io/rewrite-target: /
  name: myingress
spec:
  rules:
  - host: innoq.training.com
    http:
      paths:
      - path: /server
        backend:
          serviceName: www
          servicePort: 80
      - path: /heise
        backend:
          serviceName: wwwep
          servicePort: 80
```

Test Ingress Setup

Ingress:
`/server -> www`
`/heise -> wwwep`



Test Ingress Setup (logical)



Exercise Ingress

```
$ kubectl run echoheaders --image=gcr.io/google_containers/echoserver:1.4 --replicas=1 --port=8080
$ kubectl expose deployment echoheaders --port=80 --target-port=8080 --name=echoheaders-x
$ kubectl expose deployment echoheaders --port=80 --target-port=8080 --name=echoheaders-y
```

```
$ kubectl create -f testingress.yaml
```

```
$ kubectl describe ingress testingress
```

```
Name:                testingress
Namespace:           chris
Address:             185.19.31.109
Default backend:     default-http-backend:80 (<none>)
Rules:
  Host                Path    Backends
  ----                -
  innoq.training.com  /foo    echoheaders-x:80 (<none>)
                     /bar    echoheaders-y:80 (<none>)
```

testingress.yaml

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: testingress
spec:
  rules:
    - host: bootcamp.ch.innoq.io
      http:
        paths:
          - path: /foo
            backend:
              serviceName: echoheaders-x
              servicePort: 80
          - path: /bar
            backend:
              serviceName: echoheaders-y
              servicePort: 80
```

Exercise Ingress

```
$ curl bootcamp.ch.innoq.io/bar -H 'Host: innoq.training.com'
```

CLIENT VALUES:

client_address=10.244.98.6

command=GET

real path=/bar

query=nil

request_version=1.1

request_uri=http://innoq.training.com:8080/bar

SERVER VALUES:

server_version=nginx: 1.10.0 - lua: 10001

HEADERS RECEIVED:

accept=*

connection=close

host=innoq.training.com

user-agent=curl/7.51.0

x-forwarded-for=10.244.98.1

x-forwarded-host=innoq.training.com

x-forwarded-port=80

x-forwarded-proto=http

x-real-ip=10.244.98.1

BODY:

-no body in request-

Secrets

Secrets

- **intended to hold sensitive information, such as**
 - passwords
 - OAuth tokens
 - ssh keys
- **safer and more flexible than putting it verbatim in a pod definition or in a docker image**
- **reduces the risk of accidental exposure, during the workflow of creating, viewing, and editing pods.**
- **can be used in three ways:**
 - as files in a volume mounted on one or more of its containers,
 - as Environment Variables
 - or used by kubelet when pulling images for the pod

But...

- **In the API server secret data is stored as plaintext in etcd; therefore:**
 - limit access to etcd to admins
 - wipe/shred disks used by etcd
 - **Apps still need to protect the value of secret after reading it from the volume**
 - **Users who can create a pod that uses a secret can also see the value of that secret**
 - **No control which users of a Kubernetes cluster can access a secret**
- > Changes with RBAC and etcd encryption (K8s > 1.7.x)**

Example:

```
$ echo -n "admin" > ./username.txt
```

```
$ echo -n "1f2d1e2e67df" > ./password.txt
```

```
$ kubectl create secret generic db-user-pass --from-file=./username.txt --from-file=./password.txt  
secret "db-user-pass" created
```

```
$ kubectl describe secret db-user-pass
```

```
Name:          db-user-pass
```

```
Namespace: chris
```

```
Labels:         <none>
```

```
Annotations:    <none>
```

```
Type: Opaque
```

```
Data
```

```
====
```

```
password.txt:   12 bytes
```

```
username.txt:   5 bytes
```

Example:

```
$ kubectl get secret db-user-pass -o yaml
apiVersion: v1
data:
  password.txt: MWYyZDFlMmU2N2Rm
  username.txt: YWRtaW4=
kind: Secret
metadata:
  creationTimestamp: 2017-01-30T12:57:21Z
  name: db-user-pass
  namespace: chris
  resourceVersion: "692297"
  selfLink: /api/v1/namespaces/chris/secrets/db-user-pass
  uid: a33a0e3f-e6eb-11e6-903c-06c92a0003b6
type: Opaque
```


Exercise Secrets

1. *Create secret mysecret*
2. *Try to decode it with ... get ... -o yaml and ... base64 --decode*
3. *Review / create secretbusybox Pod and view its log*
4. *kubectl exec -it secretbusybox sh*
And cat both files in /etc/secrets
5. *BTW:*

What are all these Env Variables?
6. *Does a secret get updated?*

mysecret.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDFlMmU2N2Rm
```

Config Maps

Config Map

holds key-value pairs of configuration data

similar to Secrets, but designed to more conveniently support working with strings that do not contain sensitive information

can be used to:

- **Populate the values of environment variables**
- **Set command-line arguments in a container**
- **Populate config files in a volume**

Config Map

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: special-config
data:
  special.how: very
  special.type: charm
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: env-config
data:
  log_level: INFO
  stage: INT
```

```
apiVersion: v1
kind: Pod
metadata:
  name: dapi-test-pod
spec:
  containers:
    - name: test-container
      image: gcr.io/google_containers/busybox
      command: [ "/bin/sh", "-c", "env" ]
      env:
        - name: SPECIAL_LEVEL_KEY
          valueFrom:
            configMapKeyRef:
              name: special-config
              key: special.how
        - name: SPECIAL_TYPE_KEY
          valueFrom:
            configMapKeyRef:
              name: special-config
              key: special.type
      envFrom:
        - configMapRef:
            name: env-config
      restartPolicy: Never
```

Exercise Config Maps

1. Try to create ConfigMaps (see README in Creating-ConfigMaps)
2. Review and try the Config Map examples (3 Use Cases)
3. What happens if a config map does not exist? Or a Key/Value?

Hints:

`kubectl get po --show-all` *(shows finished PODs)*

`kubectl get configmap game-config -o yaml` *(shows CM content)*

Container Bootcamp

Kubernetes

10: Custom Resource Definitions (CRD)



Custom Resource Definition

Kubernetes comes with many built-in API objects

CRD extend Kubernetes with their own API object type in order to do custom automation (API Server creates REST Path)

Each CRD has the following:

- **metadata - Standard Kubernetes object metadata.**
- **kind - The kind of the resources described by this third party resource.**
- **description - A free text description of the resource.**
- **versions - A list of the versions of the resource.**

Example innoq-cron-tab

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: innoqcrontabs.innoq.com
spec:
  group: innoq.com
  names:
    kind: InnoqCronTab
    listKind: InnoqCronTabList
    plural: innoqcrontabs
    singular: innoqcrontab
    shortNames:
      - ict
  scope: Namespaced
  version: v1
  description: "A specification of a Pod to run on a cron
style schedule"
```

Creates a new RESTful API endpoint at

`/apis/stable.innoq.com/v1/namespaces/<namespace>/innoqcrontabs/...`

-> API endpoint URL and support for CRUD operations, and watch API

You can then manage our InnoqCronTab objects using kubectl

Example innoq-cron-tab

```
apiVersion: innoq.com/v1
kind: InnoqCronTab
metadata:
  name: my-new-cron-object
cronSpec: "* * * *"
image: innoqs-awesome-cron-image
```

InnoqCronTab instance `my-new-cron-object`

Exercise CRD

1. *Create the resource `innocron-tab.stable.innocron.com` with `resource.yaml`*
2. *Create an instance of `InnocronTab`*
3. *Use `get`, `describe` to review the crd / `InnocronTab`*

Introducing Operators

Operator

... is an application-specific controller that extends the Kubernetes API to


- **create**
- **configure**
- **manage**

instances of complex stateful applications


It builds upon the basic Kubernetes resource and controller concepts

Includes domain or application-specific knowledge to automate common tasks

Operator / Controller



Example: etcd Operator



Exercise/Demo CoreOS etcd Operator

1. Review and create the Operator (wait for startup)

2. Review and create a etcd-Cluster

Use `watch kubectl get po` to view cluster instantiation. Services?

3. Put some data in it

```
$ kubectl run --rm -i --tty fun --image quay.io/coreos/etcd --restart=Never -- /bin/sh
/ # export ETCDCTL_API=3
/ # etcdctl --endpoints http://example-etcd-cluster-client.etcd-operator:2379 put foo bar
OK
/ # etcdctl --endpoints http://example-etcd-cluster-client.etcd-operator:2379 get foo
foo
bar
(ctrl-D to exit)
```

4. Delete Cluster my-etcd-cluster