

Container Bootcamp

# Microservice Operations

INNOQ

# Operations

- **Huge challenge**
- **Need to operate 50-100 Microservices**
- **More (virtual) system than an IT department might have**

# Provisioning

# Provisioning

- **Manual deployments too much effort**
- **...and too slow**
- **Manual deployments: hard to get right**
- **...and hard to reproduce**
  
- **So: Automate!**

# Provisioning Tools

- Your favorite package manager (apt-get, yum, HELM ...)
- ...but: (operating) system dependent
- Puppet , Chef, Ansible, Salt
- Your custom solution

# Idempotent Provisioning

- **Idea: Describe desired state of system after deployment**
- **Each install run gives the same result**
- **Machine fresh -> full install**
- **Machine up to date -> nothing**

# Idempotent Provisioning

- **Problem: Complex**
- **No scripts, but declarations**
- **Problem: Might be incomplete**
- **Problem: Really idempotent?**

# Immutable Server

- **Idea: Server cannot be changed**
- **Instead: Provide a complete new server for each update**
- **Technologies like Docker enable this**




# Log Analysis

# Log Analysis


- **Log Files are easy to analyze**
- **But: Microservices are distributed**
- **...so are logs**
- **Must analyze all logs**
- **Centralized log analysis**

# ELK Stack



- **Logstash: JRuby**
- **Inputs, Parser, Outputs**
- **Elasticsearch: Java**
- **Kibana: JavaScript**

# EFK Stack



# Graylog

- **Open Source**
- **Elasticsearch for storage**
- **MongoDB for Meta data**
- **GELF format for log messages**

# More Alternatives

- **Splunk: commercial**
- **+Cloud**
- **Cloud : Loggly**
- **Sumo Logic**
- **Papertrail**

# Monitoring

# App Metrics

- **Covers production relevant technical service metrics**
- **E.g. service quality degrading without a complete service failure**
  - Pull principle, for things like:
    - Number of HTTP 500 response codes in the last 10 minutes.
    - Average DB query duration
- **Could be used for Auto-Scaling**
- **f.e. Prometheus**
  - While queries should be written by service developer



# Node Metrics

**Node Metrics are still valuable, to get an insight to per node HW utilization**

- **Memory load**
- **Processes**
- **File descriptors**
- **Network throughput**

# Business Metrics

**There might be a need for measuring business metrics.  
Depending on the stakeholder this might require a different  
technical solution**

# Prometheus

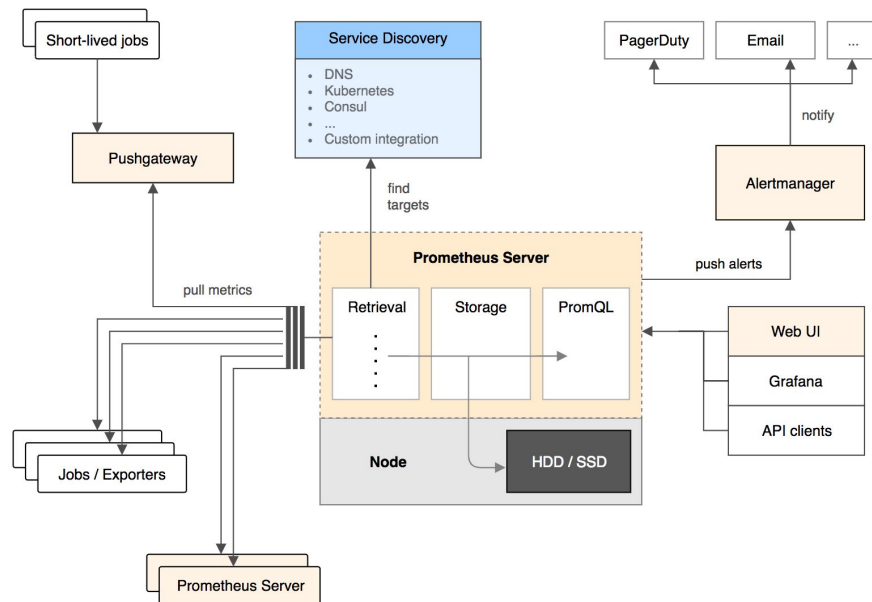
Is basically a Time Series Database

Polls data from a /metrics endpoint

Complex query language

Alerting

Visualization



# Tracing

# Tracing

- **How is a request/workflow handled throughout the system?**
- **Trace call in each Microservice**
- **E.g. log + HTTP header**

# Standardize

- **Consider to standardize**
- **Log format**
- **How logs are written**
- **How values are forwarded to monitoring**
- **Tracing**
- **Simplifies operations**
- **Individual system have no benefits**

# DevOps?

- **DevOps = Development + Operations**
- **Close collaboration**
- **Ops skill become more important for Dev**
- **But: Do you need DevOps for Microservices?**
- **No – Ops define Macro-Architecture only**

# Conclusion



# Conclusion

- **Automate Provisioning**
- **Central log processing**
- **Central monitoring**
- **Standardize!**
- **DevOps beneficial but not required**